

# Reproducing “FlashRoute: Efficient Traceroute on a Massive Scale<sup>1</sup>”

Ziqi Zhao  
Zhengdong Wang  
Hao Yin  
Liangtai Sun

December 14, 2020

---

<sup>1</sup>Yuchen Huang, Michael Rabinovich, and Rami Al-Dalky. 2020. FlashRoute: Efficient Traceroute on a Massive Scale. In ACM Internet Measurement Conference (IMC '20), October 27–29, 2020, Virtual Event, USA. ACM, New York, NY, USA, 13 pages. <https://doi.org/10.1145/3419394.3423619>

# Table of Contents

## 1 Motivation

- Traceroute
- FlashRoute

## 2 Implementation

## 3 Results

- Performance
- Impact of Some Features
- Topology Visualization

## 4 Appendix

# Table of Contents

## 1 Motivation

- Traceroute
- FlashRoute

## 2 Implementation

## 3 Results

- Performance
- Impact of Some Features
- Topology Visualization

## 4 Appendix

# Table of Contents

- 1 Motivation
  - Traceroute
  - FlashRoute
- 2 Implementation
- 3 Results
- 4 Appendix

# Traceroute

- Traditional traceroute




Figure: An example of traceroute<sup>2</sup>

<sup>2</sup>[https://major.io/wp-content/uploads/2012/06/RAS\\_Traceroute\\_Book\\_Format.pdf](https://major.io/wp-content/uploads/2012/06/RAS_Traceroute_Book_Format.pdf)

# Traceroute

- Traditional traceroute




Figure: An example of traceroute<sup>2</sup>

- What if on massive-scale networks, for example, /24 ?

<sup>2</sup>[https://major.io/wp-content/uploads/2012/06/RAS\\_Traceroute\\_Book\\_Format.pdf](https://major.io/wp-content/uploads/2012/06/RAS_Traceroute_Book_Format.pdf)

# What if on massive-scale networks?

Some critical issues...

# What if on massive-scale networks?

Some critical issues...

- Too slow
  - ▶ send & receive synchronously
  - ▶ often blocked due to unresponsive servers



# What if on massive-scale networks?

Some critical issues...

- Too slow
  - ▶ send & receive synchronously
  - ▶ often blocked due to unresponsive servers
- Too aggressive
  - ▶ routers must respond a huge number of times
  - ▶ may be mistaken for a DoS attack

# What if on massive-scale networks?

Some critical issues...

- Too slow
  - ▶ send & receive synchronously
  - ▶ often blocked due to unresponsive servers
- Too aggressive
  - ▶ routers must respond a huge number of times
  - ▶ may be mistaken for a DoS attack

How can we solve these? **FlashRoute**.

# Table of Contents

- 1 Motivation
  - Traceroute
  - FlashRoute
- 2 Implementation
- 3 Results
- 4 Appendix

# FlashRoute Methodologies

- Asynchronous sending and receiving
  - Encode probing context into the IP and UDP headers

| IP Bit Offset  | 0-3                 | 4-7           | 8-13     | 14-15           | 16-18                   | 19-31 |
|----------------|---------------------|---------------|----------|-----------------|-------------------------|-------|
| 0              | Version             | Header Length | DSCP     | ECN             | Total Length            |       |
| 32             | Identification      |               |          | Flags           | Fragment Offset         |       |
| 64             | Time to Live        |               | Protocol | Header Checksum |                         |       |
| 96             | Source Address      |               |          |                 |                         |       |
| 128            | Destination Address |               |          |                 |                         |       |
| UDP Bit Offset | 0-15                |               |          |                 | 16-31                   |       |
| 160            | Source Port Number  |               |          |                 | Destination Port Number |       |
| 192            | Length              |               |          |                 | Checksum                |       |
| 224            | Payload             |               |          |                 |                         |       |

Handwritten annotations on the table:

- Red arrow from ECN (14-15) to Total Length (19-31) labeled "Timestamp".
- Red arrow from Identification (32-47) to Total Length (19-31) labeled "6".
- Red arrow from Identification (32-47) to Time to Live (64-79) labeled "10".
- Red arrow from Protocol (80-95) to Time to Live (64-79) labeled "TTL".
- Red arrow from Protocol (80-95) to Header Checksum (96-111) labeled "phase".
- Red arrow from Source Port Number (160-175) to Checksum (192-207) labeled "Checksum of Src IP".

Figure: Packet encoding

# FlashRoute Methodologies (Cont.)

- Doubletree algorithm
  - ▶ Split point or TTL
  - ▶ Prefer backward probing and reduce redundancy to a minimum




Figure: An ideal topology<sup>3</sup>

<sup>3</sup>Figure 1 in the paper

# FlashRoute Methodologies (Cont.)

- Preprobing
  - ▶ Split TTL too large → mistaken for a DoS attack
  - ▶ Split TTL too small → miss probes-saving opportunities that backward probing exploits

# FlashRoute Methodologies (Cont.)

- Preprobing

- ▶ Split TTL too large → mistaken for a DoS attack
- ▶ Split TTL too small → miss probes-saving opportunities that backward probing exploits
- ▶ One-probe hop-distance measurement
- ▶ Provide an estimate of distances to targets!

# Table of Contents

## 1 Motivation

- Traceroute
- FlashRoute

## 2 Implementation

## 3 Results

- Performance
- Impact of Some Features
- Topology Visualization

## 4 Appendix



# *flashroute.rs*: A Rust Implementation

Why Rust<sup>4</sup>?

- A modern system programming language

---

<sup>4</sup><https://rust-lang.org>

# flashroute.rs: A Rust Implementation


Why Rust<sup>4</sup>?

- A modern system programming language
- Performance: blazingly fast and memory-efficient
- Reliability: **ownership model** guarantees memory-safety and thread-safety
- Productivity: great community provides a lot of high-performance 3rd-party libraries


---

<sup>4</sup><https://rust-lang.org>

# Redesign in the view of Task - Overview




# Redesign in the view of Task - Overview




- Lightweight, coroutine-level, automatic task scheduling.
- Make full use of the performance of modern multi-core processors.

# Redesign in the view of Task - Thread-safety



- Mutex or rwlock *FREE*. All inter-task communications are achieved through message channels or atomic operations.

# Redesign in the view of Task - Thread-safety



- Mutex or rwlock *FREE*. All inter-task communications are achieved through message channels or atomic operations.
- *Absolute* thread-safety thanks to Rust's ownership model.
- Much lower overhead.

# Swiss Table Based Data Structure

- Originally targets are organized as an array-based implicit linked list
  - ▶ which is enough for the purpose of only scanning /24 internet...
  - ▶ but with poor scalability

# Swiss Table Based Data Structure

- Originally targets are organized as an array-based implicit linked list
  - ▶ which is enough for the purpose of only scanning /24 internet...
  - ▶ but with poor scalability
- Swiss table: a high-performance hash table
- Advantages
  - ▶ allow customizing probing grain
  - ▶ probe partial internet with much lower memory usage



# Human-readable Output

- *flashroute.rs* provides the ability to probe on any subnet

# Human-readable Output

- *flashroute.rs* provides the ability to probe on any subnet
- Propose a method to assemble traceroute results into a graph

---

```
graph {  
  0 [ label = "192.168.1.177" ]  
  1 [ label = "115.159.1.131" ]  
  2 [ label = "9.31.253.89" ]  
  ...  
  2 -- 1 [ label = "1" ]  
  12 -- 11 [ label = "5" ]  
  7 -- 13 [ label = "4" ]  
  ...  
}
```

---

# Human-readable Output

- *flashroute.rs* provides the ability to probe on any subnet
- Propose a method to assemble traceroute results into a graph

---

```
graph {  
  0 [ label = "192.168.1.177" ]  
  1 [ label = "115.159.1.131" ]  
  2 [ label = "9.31.253.89" ]  
  ...  
  2 -- 1 [ label = "1" ]  
  12 -- 11 [ label = "5" ]  
  7 -- 13 [ label = "4" ]  
  ...  
}
```

---

- Generate visualization on small-scale topology

# Table of Contents

## 1 Motivation

- Traceroute
- FlashRoute

## 2 Implementation

## 3 Results

- Performance
- Impact of Some Features
- Topology Visualization

## 4 Appendix

# Table of Contents

1 Motivation

2 Implementation

3 Results

- Performance
- Impact of Some Features
- Topology Visualization

4 Appendix

# Overview of *flashroute.rs*

Our goals:

- Enhance FlashRoute in **Rust**
  - ▶ originally implemented in C++ using 2.4k LoC<sup>5</sup>
  - ▶ expectations
    - ★ fewer codes
    - ★ even better performance
    - ★ asynchronous operations instead of manually-handled mutex

---

<sup>5</sup><https://github.com/lambdahuang/FlashRoute/>

<sup>6</sup>Tested on AMD EPYC 7B12 (8) @ 2.25 GHz

# Overview of *flashroute.rs*

Our goals:

- Enhance FlashRoute in **Rust**
  - ▶ originally implemented in C++ using 2.4k LoC<sup>5</sup>
  - ▶ expectations
    - ★ fewer codes
    - ★ even better performance
    - ★ asynchronous operations instead of manually-handled mutex

Our results:

- *flashroute.rs* succeeds to probe 800,000 interfaces as FlashRoute does

---

<sup>5</sup><https://github.com/lambdahuang/FlashRoute/>

<sup>6</sup>Tested on AMD EPYC 7B12 (8) @ 2.25 GHz

# Overview of *flashroute.rs*

Our goals:

- Enhance FlashRoute in **Rust**
  - ▶ originally implemented in C++ using 2.4k LoC<sup>5</sup>
  - ▶ expectations
    - ★ fewer codes
    - ★ even better performance
    - ★ asynchronous operations instead of manually-handled mutex

Our results:

- *flashroute.rs* succeeds to probe 800,000 interfaces as FlashRoute does
- *flashroute.rs* implements almost all features in only **1.2k LoC**

---

<sup>5</sup><https://github.com/lambdahuang/FlashRoute/>

<sup>6</sup>Tested on AMD EPYC 7B12 (8) @ 2.25 GHz



# Overview of *flashroute.rs*

Our goals:

- Enhance FlashRoute in **Rust**
  - ▶ originally implemented in C++ using 2.4k LoC<sup>5</sup>
  - ▶ expectations
    - ★ fewer codes
    - ★ even better performance
    - ★ asynchronous operations instead of manually-handled mutex

Our results:

- *flashroute.rs* succeeds to probe 800,000 interfaces as FlashRoute does
- *flashroute.rs* implements almost all features in only **1.2k LoC**
- *flashroute.rs* gains probing performance improvement of up to **40%**<sup>6</sup>

---

<sup>5</sup><https://github.com/lambdahuang/FlashRoute/>

<sup>6</sup>Tested on AMD EPYC 7B12 (8) @ 2.25 GHz

# Performance of Single Target Probing

Results of probing single target 115.159.1.233 from 59.78.X.X:

| Tool                 | Probe Time     |
|----------------------|----------------|
| <i>flashroute.rs</i> | <b>0:08.34</b> |
| FlashRoute           | 0:15.19        |
| Scamper              | 0:50.55        |
| tracert              | 0:55.10        |
| Yarrp                | 1:01.51        |

## Performance at Probing Rate of 100 Kpps

Results with different split TTL, at probing rate of 100 Kpps:

| Tool                     | Interfaces | Probes      | Scan Time | Probing Rate |
|--------------------------|------------|-------------|-----------|--------------|
| <i>flashroute.rs</i> -16 | 822,527    | 180,282,633 | 30:20     | 99,056       |
| FlashRoute-16            | 824,923    | 162,546,005 | 27:24     | 98,872       |
| FlashRoute-16*           | 812,403    | 97,807,092  | 17:17     | 94,317       |
| <i>flashroute.rs</i> -32 | 818,299    | 376,085,798 | 1:02:38   | 100,076      |
| FlashRoute-32            | 820,745    | 342,091,698 | 57:07     | 99,822       |
| FlashRoute-32*           | 807,588    | 159,185,459 | 27:32     | 96,359       |
| Scamper-16               | 805,970    | 217,940,057 | 6:43:53   | 8,993        |

The results marked with \* are collected from the paper, which performs much better since it preprobes the targets according to the Census Hitlist dataset<sup>7</sup>, instead of a randomly-generated one.

<sup>7</sup>USC/LANDER project. Jan 30 2019. Internet Addresses Hitlist Dataset. USCLANDER/internet\_address\_hitlist\_it84w-20190130/rev10351.

# Performance at Maximum Probing Rate

Results with different split TTL, at maximum probing rate:

| Tool                     | Interfaces | Probes      | Scan Time    | Probing Rate   |
|--------------------------|------------|-------------|--------------|----------------|
| <i>flashroute.rs</i> -8  | 840,437    | 136,655,629 | <b>9:09</b>  | <b>248,917</b> |
| FlashRoute-8             | 847,343    | 127,698,272 | 12:39        | 168,245        |
| <i>flashroute.rs</i> -16 | 834,601    | 180,874,492 | <b>12:01</b> | <b>250,866</b> |
| FlashRoute-16            | 837,672    | 163,091,546 | 15:41        | 173,317        |
| <i>flashroute.rs</i> -32 | 830,889    | 377,175,922 | <b>22:52</b> | <b>274,910</b> |
| FlashRoute-32            | 835,343    | 342,231,468 | 30:29        | 187,113        |
| Scamper-16               | 805,970    | 217,940,057 | 6:43:53      | 8,993          |

# Performance for Different CPU Configurations

| Tool                 | CPUs | Interfaces | Probes      | Scan Time    | Rate           |
|----------------------|------|------------|-------------|--------------|----------------|
| <i>flashroute.rs</i> | 4    | 803,248    | 203,197,113 | <b>21:00</b> | <b>161,267</b> |
| FlashRoute           | 4    | 804,711    | 185,251,917 | 24:54        | 123,997        |
| <i>flashroute.rs</i> | 8    | 834,601    | 180,874,492 | <b>12:01</b> | <b>250,866</b> |
| FlashRoute           | 8    | 837,672    | 163,091,546 | 15:41        | 173,317        |

Results are tested on

- 4 CPUs: Intel Xeon (4) @ 2.0 GHz
- 8 CPUs: AMD EPYC 7B12 (8) @ 2.25 GHz

# Table of Contents

1 Motivation

2 Implementation

3 Results

- Performance
- **Impact of Some Features**
- Topology Visualization

4 Appendix

# Impact of Redundancy Removal during Backward Probing

| Tool                    | RR <sup>8</sup> | Interfaces | Probes             | Received          | Time         |
|-------------------------|-----------------|------------|--------------------|-------------------|--------------|
| <i>flashroute.rs-8</i>  | On              | 840,437    | <b>136,655,629</b> | <b>26,746,697</b> | <b>9:09</b>  |
| <i>flashroute.rs-8</i>  | Off             | 838,082    | 204,739,320        | 60,359,387        | 12:54        |
| <i>flashroute.rs-16</i> | On              | 834,601    | <b>180,874,492</b> | <b>14,491,139</b> | <b>12:01</b> |
| <i>flashroute.rs-16</i> | Off             | 845,655    | 266,010,541        | 60,904,277        | 18:39        |

<sup>8</sup>Redundancy Removal

# Impact of Gap-limit




Figure: Interfaces and Scan Time under Different Gap-limit



# Table of Contents

1 Motivation

2 Implementation

3 Results

- Performance
- Impact of Some Features
- **Topology Visualization**

4 Appendix

## Visualization




Figure: Topology of some hosts @115.159.1.0/24

# Table of Contents

## 1 Motivation

- Traceroute
- FlashRoute

## 2 Implementation

## 3 Results

- Performance
- Impact of Some Features
- Topology Visualization

## 4 Appendix

# Contribution Matrix

| Item            | Z. Zhao | Z. Wang | H. Yin | L. Sun |
|-----------------|---------|---------|--------|--------|
| Coding          | •       |         |        |        |
| Code Review     |         | •       | •      | •      |
| Code Test       | •       | •       |        |        |
| Experiment      | •       | •       |        |        |
| Data Processing | •       | •       | •      | •      |
| Slides          | •       | •       | •      | •      |
| Report          | •       | •       | •      | •      |
| Google Cloud    | •       |         |        |        |

# Sources

All sources of *flashroute.rs* can be found at:  
<https://github.com/BugenZhao/flashroute.rs>

# Thank You